

# R code optimization and profiling

Alex Shlemov

✉ [a.shlemov@spbu.ru](mailto:a.shlemov@spbu.ru)

🌐 eodus

Saint Petersburg State University, Russia  
Faculty of Mathematics and Mechanics  
Department of Statistical Modelling

March 30, 2015



- 1 Introduction
- 2 The R way, patterns and antipatterns
- 3 “Something for nothing”
- 4 Some low-level optimization techniques
- 5 Code profiling
- 6 Conclusions and further studying


**R** is very good:

- High-level language with a lot of useful features
- Friendly, smooth learning curve
- Much cool out-of-box stuff
- Many packages
- Broad community
- Open and free
- Cross-platform (really cross-platform!)
- Easily extendable
- ...

**R** is very good:

- High-level language with a lot of useful features
- Friendly, smooth learning curve
- Much cool out-of-box stuff
- Many packages
- Broad community
- Open and free
- Cross-platform (really cross-platform!)
- Easily extendable
- ...

...but sometimes it's very slow!



# Example

C++ (via 'Rcpp'):

```
#include <Rcpp.h>
#include <cmath>

// [[Rcpp::export]]
double fcpp(const int n) {
    double s = 0.;
    for (int i = 1; i <= n; ++i)
        s += sin(i);

    return s;
}

fcpp(1e6)

## [1] -0.1171095

system.time(fcpp(1e6))

##      user  system elapsed
##  0.055   0.000   0.055
```

R (Revolution):

```
f <- function(n) {
    s <- 0
    for (i in 1:n)
        s <- s + sin(i)

    s
}

f(1e6)

## [1] -0.1171095

system.time(f(1e6))

##      user  system elapsed
##  0.521   0.004   0.525
```

# 'microbenchmark' timing package

`system.time()` works, but for fast calls it may be unprecise  
microbenchmark from the package 'microbenchmark' is better alternative:

```
library(microbenchmark)
```

```
microbenchmark(f(1e6), fcpp(1e6), times = 25)
```

```
## Unit: milliseconds
```

```
##      expr      min       lq      mean    median      uq      max neval
##  f(1e+06) 518.58249 526.68754 548.77125 531.17251 551.81352 721.70462    25
## fcpp(1e+06) 54.78589 54.90963 55.35176 55.00016 55.25229 60.18847    25
```

`benchmark()` from 'rbenchmark' is quite good too:

```
library(rbenchmark)
```

```
benchmark(f(1e6), fcpp(1e6), replications = 25)
```

```
##      test replications elapsed relative user.self sys.self user.child sys.child
## 1    f(1e+06)         25  13.252    9.659    13.238    0.012         0         0
## 2 fcpp(1e+06)         25   1.372    1.000    1.372    0.000         0         0
```

## Why is R code so slow?.. Main three causes:

- Often **R** code is poor written (nothing to add)
- **R** implementation is not optimized
- **R** is extremely *dynamic* language (all types are polymorphic, we can write to every environment in any time, etc)

See <http://adv-r.had.co.nz/Performance.html> for comprehensive explanation

## How to speedup R code?

- Write better code following the **R** way
- Use optimized **R** implementations
- (If we really need) Reduce “dynamism” and/or use more low-level approaches (give up polymorphism, use more low-level functions or more low-level languages and so on)

*“Programmers waste enormous amounts of time thinking about, or worrying about, the speed of noncritical parts of their programs, and these attempts at efficiency actually have a strong negative impact when debugging and maintenance are considered”*

— Donald Knuth



- 1 Introduction
- 2 The R way, patterns and antipatterns
- 3 “Something for nothing”
- 4 Some low-level optimization techniques
- 5 Code profiling
- 6 Conclusions and further studying

- 1 Introduction
- 2 The R way, patterns and antipatterns
- 3 “Something for nothing”
- 4 Some low-level optimization techniques
- 5 Code profiling
- 6 Conclusions and further studying

*Loops are ugly and slow. Vectorized operations are elegant and fast*

Let's try:

```
library(microbenchmark)
fvec <- function(n) {
  sum(sin(1:n))
}
```

```
microbenchmark(f(1e5), fvec(1e5), fcpp(1e5), times = 10)
```

```
## Unit: milliseconds
```

##	expr	min	lq	mean	median	uq	max	neval
##	f(1e+05)	49.943049	51.378679	55.422791	51.796625	52.807415	89.197145	10
##	fvec(1e+05)	6.060916	6.245789	6.428412	6.517029	6.642355	6.682902	10
##	fcpp(1e+05)	5.506135	5.517833	5.549802	5.553273	5.568905	5.608258	10

Not bad!

# Vectorized functions

Use:

- vectorized mathematical functions like `sin()`, `log()`, etc;
- `ifelse()` as “vectorized if”;
- `colSums()` instead of `apply(X, 2, sum)`, `rowMeans()` instead of `apply(X, 1, mean)` and so on;
- `max.col(X)` rather than `apply(X, 1, which.max)`;
- vectorized subscripting, e.g. `x[is.na(x)] <- 0` or `x <- x[x %in% S]`;
- vectorized subscripting for matrices and arrays like `m[1:10, ]`, `m[cbind(1:10, 10:1)]` and `m[1:10]`

`match()` and `%in%` are also vectorized.

Be aware of some vectorized functions like `diff()` and `cumsum()`

Combine! Write your own vectorized functions

# \*apply() and their friends

\*apply(), vectorize(), replicate(), aggregate(), etc don't provide any real vectorization!

They are less efficient than vectorized functions but usually much better than ordinary loops because of it's easy to parallelize them

Package 'foreach' is good alternative for loops too (by the same reason)

Be aware of tapply() (and by()). tapply() is *grouping* apply (like SQL' GROUP BY):

```
h <- data.frame(sex = sample(c("male", "female"), size = 10, replace = TRUE),
               height = rnorm(10, mean = 170, sd = 10))
tapply(h$height, h$sex, mean)
```

```
##   female    male
## 165.8786 168.7190
```

rapply() is also sometimes useful

# Memory preallocation

**R** vectors are not like C++ `std::vector`'s, **R** don't reserve any space for “appending”. Each resizing invokes reallocation with  $\mathcal{O}(l)$  complexity. Thus, such code:

```
x <- c()
for (i in 1:n) {
  x <- c(x, f(i))
}
```

has  $\mathcal{O}(n^2)$  time complexity! Avoid growing objects!  
Always try to use vectorization. If it's impossible:

```
# x <- numeric(n) # NA-filling is better because of preventing some errors
x <- rep(NA_real_, n) # Preallocation, suppose that all results are numeric
for (i in seq_len(n)) { # seq_len is a bit faster
  x[i] <- f(i)
}
```

`*apply()` function family almost always provides better alternative, e.g.:

```
x <- sapply(1:n, f) # Shorter
x <- vapply(seq_len(n), f, f(1)) # Little bit faster because of more proper preallocation
```

# If you really need growing object...

...Use preallocation like in `std::vector`

Bad:

```
res <- c()
i <- 1
while(TRUE) {
  cur <- f(i)
  if (g(cur)) { # Some stop condition
    break
  } else {
    res <- c(res, cur)
  }
  i <- i + 1
}
```

Much better:

```
res <- c()
i <- 1
while(TRUE) {
  cur <- f(i)
  if (g(cur)) { # Some stop condition
    break
  } else {
    if (length(res) < i)
      res[2 * i] <- NA_real_
    res[i] <- cur
  }
  i <- i + 1
}
res <- res[seq_len(i - 1)] # Truncate
```

# Prefer in-place operations

In-place operations are often much faster:

```
squish_ife <- function(x, a, b) {  
  ifelse(x <= a, a, ifelse(x >= b, b, x))  
}  
squish_p <- function(x, a, b) {  
  pmax(pmin(x, b), a)  
}  
squish_in_place <- function(x, a, b) {  
  x[x <= a] <- a  
  x[x >= b] <- b  
  x  
}
```

```
x <- runif(100, -1.5, 1.5)  
microbenchmark(squish_ife(x, -1, 1), squish_p(x, -1, 1), squish_in_place(x, -1, 1))
```

```
## Unit: microseconds
```

##	expr	min	lq	mean	median	uq	max	neval
##	squish_ife(x, -1, 1)	97.565	99.0395	101.43645	100.2890	101.7715	125.794	100
##	squish_p(x, -1, 1)	28.818	31.1260	36.04667	32.9915	33.8885	269.673	100
##	squish_in_place(x, -1, 1)	11.348	12.3850	13.49622	13.2710	13.6735	34.457	100



# Chunking for large structures

Remember our “sum of sines”. Let  $n$  be very big:

```
fvec <- function(n) {  
  sum(sin(seq_len(n))) # <= Huge vector constructed here  
}
```

```
f(1e10) # CRASH!
```

```
## Error: cannot allocate vector of size 74.5 Gb
```

Try to devectorize (chunk):

```
fchunk <- function(n, chunk = 1e7) {  
  sum(vapply(seq(from = 1, by = chunk, length.out = ceiling(n / chunk)),  
    function(i) sum(sin(seq(i, min(i + chunk - 1, n)))),  
    0.))  
}  
fchunk(1e10)
```

```
## [1] -0.1276265
```

# Use proper methods, functions, data structures and packages

There are many **R** packages...  
...and some of them are better than others

Find proper packages for your problem!

Try to use, e.g:

- `'data.table'` and `'sqldf'` for big dataframes;
- `'zoo'`, `'xts'` for time series;
- `'mboost'`, `'flare'` for linear regression;
- `'optimx'` for optimization;
- `'cluster'` for clustering;
- `'Matrix'` for large sparse matrices;
- `'svd'` for fast truncated SVD (used e.g. for PCA and related methods);
- `'FFTW'` for Fourier transform (rather than slow `fft()`);
- ...

*K-means* clustering with multistart (single-thread and parallel versions):

```
best.kmeans <- function(x, centers, N) {  
  cls <- lapply(seq_len(N),  
               function(i) kmeans(x, centers = centers))  
  cls[[which.min(sapply(cls, function(x) x$tot.withinss))]]  
}
```

```
best.kmeans.par <- function(x, centers, N, cluster) {  
  cls <- parLapplyLB(cluster,  
                    seq_len(N),  
                    function(i) kmeans(x, centers = centers))  
  cls[[which.min(sapply(cls, function(x) x$tot.withinss))]]  
}
```

Package 'foreach' provides more friendly interface

# Parallelize

```
library(parallel)
cores <- detectCores()
cores

## [1] 4

cluster <- makePSOCKcluster(cores)

USA.scaled <- scale(USArrests)

clusterExport(cluster, "USA.scaled")

microbenchmark(singe = best.kmeans(USA.scaled, 5, 1000),
               parallel = best.kmeans.par(USA.scaled, 5, 1000, cluster), times = 10)

## Unit: milliseconds
##      expr      min       lq      mean   median       uq      max  neval
##    singe 487.1285 498.6937 512.6206 511.520 529.3923 536.3230    10
## parallel 234.5809 237.4995 263.4459 242.605 288.8327 353.0925    10

stopCluster(cluster)
```

- 1 Introduction
- 2 The R way, patterns and antipatterns
- 3 “Something for nothing”**
- 4 Some low-level optimization techniques
- 5 Code profiling
- 6 Conclusions and further studying

# Byte-compilation and JIT

**R** is *interpreter*. Originally **R** considers a function as a raw text and parses it on each call. Byte-compilation can reduce time costs for expression parsing:

```
f <- function(n) {  
  s <- 0  
  for (i in 1:n) s <- s + sin(i)  
  
  s  
}
```

```
library(compiler)  
fcomp <- cmpfun(f)
```

```
microbenchmark(f(1e5), fcomp(1e5), times = 100)
```

```
## Unit: milliseconds
```

##	expr	min	lq	mean	median	uq	max	neval
##	f(1e+05)	47.44344	48.63555	50.74597	48.86885	50.06414	97.74632	100
##	fcomp(1e+05)	25.86107	26.31081	27.87162	27.29993	27.91099	40.95275	100

JIT means *Just-In-Time* [compilation]. If JIT is enabled, all functions are automatically compiled before their first use

```
enableJIT(3) # 3 means the most aggressive compilation
```

## Pro:

- Something for nothing, easy to try, easy to give up
- Sometimes significantly speeds up complicated code

## Contra:

- Byte-compiled code cannot be profiled
- Much speedup only for loops; almost no speedup for vectorized code
- Sometimes slows up code

# Alternative math libraries

blas and LAPACK are common used libraries for matrix operations and linear algebra

ATLAS, OpenBLAS, Intel MKL, AMD ACML provide more efficient implementation with the same interface

On my Intel I3 laptop:

```
mx <- matrix(rnorm(10002), 1000, 1000)
```

*#with default BLAS and LAPACK*

```
system.time(svd(mx))
```

```
##      user  system elapsed
```

```
##  4.734    0.016    4.751
```

*# with OpenBLAS and ATLAS LAPACK*

```
system.time(svd(mx))
```

```
##      user  system elapsed
```

```
##  3.391    0.348    1.942
```



## Pro:

- Speedup for all linear algebra procedures ( $Ax = b$ , PCA, LM, GLM, ...)
- Easy to install, easy to uninstall (on Linux)
- Speedups not only R code (also Python 'numpy', GNU Octave, ...)

## Contra:

- Significant speedup only for large linear algebra problems
- Possible bugs and incompatibilities
- Non-trivial to install on Windows
- Intel MKL and AMD ACML are the most efficient but both are commercial

*“Revolution **R** Enterprise is the fastest, most cost effective enterprise-class big data big analytics platform available today”*

Revolution R is optimized R distribution, which already includes Intel MKL

Even faster than OpenBLAS:

```
mx <- matrix(rnorm(1000^2), 1000, 1000)
```

```
#with Revolution (Intel MKL) BLAS and LAPACK
```

```
system.time(svd(mx))
```

```
##      user  system elapsed
```

```
##  2.425    0.012    1.365
```

```
# with OpenBLAS and ATLAS LAPACK
```

```
system.time(svd(mx))
```

```
##      user  system elapsed
```

```
##  3.391    0.348    1.942
```

## Pro:

- Open edition is free
- Easy to install (anywhere). Just execute installer
- Already provides Intel MKL (for R only) for all supported platforms
- Some common **R** packages are optimized too

## Contra:

- Commercial-based; less community support; possible copyright/license issues
- Possible bugs and incompatibilities (e.g. with 'svd', 'devtools')
- Some packages may have old versions

- 1 Introduction
- 2 The R way, patterns and antipatterns
- 3 “Something for nothing”
- 4 Some low-level optimization techniques**
- 5 Code profiling
- 6 Conclusions and further studying

# Get rid of S3/S4 method dispatching

S3/S4 method dispatching (polymorphism) is very slow (R isn't C++ or Java)

E.g., function `mean()` is generic (S3 polymorphic function):

```
methods(mean)

## [1] mean.Date      mean.default    mean.difftime  mean.POSIXct   mean.POSIXlt

x <- runif(1e2)

microbenchmark(mean(x), mean.default(x))

## Unit: microseconds
##          expr    min      lq    mean median      uq    max neval
##      mean(x) 6.942  7.1705  8.32311  7.3060  7.6225 60.489   100
## mean.default(x) 1.929  2.1195  2.33768  2.1925  2.3645 10.059   100
```

Thus, for S3 specify function directly, for S4 use `findMethod()` to find the method and cache it into variable

# Work on low-level

**R** is friendly and its functions contains a lot of checks and coercions. If we exclude it, we can significantly speedup our code

```
findInterval
## function (x, vec, rightmost.closed = FALSE, all.inside = FALSE)
## {
##   if (anyNA(vec))
##     stop("'vec' contains NAs")
##   if (is.unsorted(vec))
##     stop("'vec' must be sorted non-decreasingly")
##   .Internal(findInterval(as.double(vec), as.double(x), rightmost.closed,
##     all.inside))
## }
```

We can just write:

```
findInterval.fast <- function (x, vec, rightmost.closed = FALSE, all.inside = FALSE)
{
  .Internal(findInterval(as.double(vec), as.double(x), rightmost.closed,
    all.inside))
}
```

# Work on low-level

For `mean()` :

```
mean.default
```

```
## function (x, trim = 0, na.rm = FALSE, ...)
## {
##
##   if (!is.numeric(x) && !is.complex(x) && !is.logical(x)) {
##     warning("argument is not numeric or logical: returning NA")
##     return(NA_real_)
##   }
## ...
```

```
x <- runif(1e2)
```

```
microbenchmark(mean(x), mean.default(x), sum(x) / length(x), .Internal(mean(x)))
```

```
## Unit: nanoseconds
```

```
##           expr   min      lq      mean  median      uq      max  neval
##      mean(x) 7369 8495.5 12142.92   9157 10382.5 215629   100
## mean.default(x) 1967 2614.5  3278.45   2930  3353.0  27465   100
##  sum(x)/length(x) 692  986.5  1895.38   1196  1742.5  37745   100
## .Internal(mean(x)) 420  459.5   614.92    524   623.0   5811   100
```

# Call other code from R

Most simple:

You can combine R with other languages (using files, pipes, sockets, etc)

Most effective:

Also you can write your own package using Fortran, C or C++

Most comfortable:

use packages 'rJava', 'rPython' and especially 'Rcpp'



Avoid premature optimization (“root of all evil”) and “optimization for optimization”.  
Optimize only critical parts of code. Use profiler to find them

Optimized code may be erroneous. Verify optimized code by unit tests

Estimate actual speedup by timing. Don't use optimization which makes code much more complicated but not significantly faster

- 1 Introduction
- 2 The R way, patterns and antipatterns
- 3 “Something for nothing”
- 4 Some low-level optimization techniques
- 5 Code profiling**
- 6 Conclusions and further studying

*Code profiling* is dynamic program analysis that measures space or time complexity of a program

**Idea:** Run program and evaluate time of each call

Functions and packages:

- `Rprof()`, `summaryRprof()` — standard, out-of-box
- `profr` from package 'profr' (from CRAN)
- `lineprof()` from package 'lineprof' (from Hadley's GitHub)

```
install.packages("devtools")  
library(devtools)  
install_github("hadley/lineprof")  
library(lineprof)
```

# Profiling example

```
fchunk <- function(n, chunk = 1e7) {  
  sum(vapply(seq(from = 1, by = chunk, length.out = ceiling(n / chunk)),  
            function(i) sum(sin(seq(i, min(i + chunk - 1, n)))),  
            0.))  
}
```

```
Rprof(interval = 0.001)  
fchunk(10^8)
```

```
## [1] 1.713649
```

```
Rprof(NULL)  
head(summaryRprof())$by.self, 10)
```

```
##      self.time self.pct total.time total.pct  
## "sin"      1.565   94.73      1.565   94.73  
## ":"       0.050    3.03      0.050    3.03  
## "sum"      0.037    2.24      0.037    2.24
```

# Profiling example

```
head(summaryRprof())$by.total, 15)
```

##	total.time	total.pct	self.time	self.pct
## "<Anonymous>"	1.707	100.00	0	0
## "block_exec"	1.707	100.00	0	0
## "call_block"	1.707	100.00	0	0
## "evaluate"	1.707	100.00	0	0
## "evaluate_call"	1.707	100.00	0	0
## "in_dir"	1.707	100.00	0	0
## "process_file"	1.707	100.00	0	0
## "process_group"	1.707	100.00	0	0
## "process_group.block"	1.707	100.00	0	0
## "withCallingHandlers"	1.707	100.00	0	0
## "doTryCatch"	1.706	99.94	0	0
## "eval"	1.706	99.94	0	0
## "fchunk"	1.706	99.94	0	0
## "FUN"	1.706	99.94	0	0
## "handle"	1.706	99.94	0	0

# Profiling limitations

- Profiling does not extend to C code. You can see if your **R** code calls C/C++ code but not what functions are called inside of your C/C++ code
- Similarly, you can't see what's going on inside primitive functions or byte code compiled code
- If you're doing a lot of functional programming with anonymous functions, it can be hard to figure out exactly which function is being called. The easiest way to work around this is to name your functions
- Some calls are too fast to be traced by profiler
- Profiling may influence code performance

- 1 Introduction
- 2 The R way, patterns and antipatterns
- 3 “Something for nothing”
- 4 Some low-level optimization techniques
- 5 Code profiling
- 6 Conclusions and further studying**

- Use high-level optimization, avoid low-level one
- Avoid premature optimization, use profiler to find bottlenecks
- 'Rcpp' is good tool for low-level optimization
- Use timing for estimation of obtained speedup
- Use unit tests for verification of optimized code
- Some **R** distributions are faster (but may be incompatible)



## Out-of-view subjects:

- 'Rcpp', 'rPython', 'rJava'
- Writing own packages
- Memory tracing and profiling
- Profiling of compiled code
- Alternative **R** implementations ('pqR', 'Renjin', 'FastR', 'Riposte')

## Useful packages:

- 'sqldf' and 'data.table' for big dataframes
- 'testthat' for unit tests
- 'foreach' for effective loops paralleling

## Useful links:

- [Advanced R](#) by Hadley Wickham
- [The R Inferno](#) by Patrick Burns
- [R Packages](#) by Hadley Wickham
- [Writing R Extensions](#)

# Thanks for your attention!

# Thanks for your attention!

